

---

# Foundations of the C++ Concurrency Memory Model

**John Mellor-Crummey**

**Department of Computer Science  
Rice University**

**`johnmc@rice.edu`**

# Before C++ Memory Model

---

- **Prior practice**

- threaded programming within a shared address space, e.g. Pthreads
- implementations prohibit reordering memory operations with respect to synchronization operations
  - treat synchronization operations as opaque

- **Problem**

- C and C++ are single threaded languages using thread libraries
  - language is unaware of threads
- compilers also thread unaware
  - optimize programs for a single thread
  - problem: may perform optimizations that violate intended meaning of multithreaded programs

# A Familiar Example

---

Initially  $X=Y=0$

<u>T1</u>	<u>T2</u>
$R1 = X$	$R2 = Y$
$\text{If}(R1 == 1)$	$\text{If}(R2 == 1)$
$Y = 1$	$X = 1$

Is  $R1=R2=1$  allowed?

Without precise memory model definition: yes

How?

T1 T2 speculate values of X and Y respectively as 1  
validating other's speculation

# Structure Fields and Races

---

```
struct s {char a; char b} x;
```

Thread 1:  
x.a = 1;

Thread 2:  
x.b = 1;

Thread 1 not same as:  
struct s tmp = x;  
tmp.a = 1;  
x = tmp;

The point: Compilers must be aware of threads

# Register Promotion

---

## Original Code

```
for (...) {  
    ...  
    if (mt)  
        pthread_mutex_lock(...);  
    x = ... x ...  
    if (mt)  
        pthread_mutex_unlock(...);  
}
```

## Optimized Code

```
r = x;  
for (...) {  
    ...  
    if (mt) {  
        x = r;  
        pthread_mutex_lock(...);  
        r = x;  
    }  
    r = ... r ...  
    if (mt) {  
        x = r;  
        pthread_mutex_unlock(...);  
        r = x;  
    }  
}  
x = r;
```

# Why a C++ Memory Model

---

- **Problem: Hard for programmers to reason about correctness**
- **Without precise semantics, hard to reason if compiler will violate semantics**
- **Compiler transformations could introduce data races without violating language specification**
  - resulting executions could yield unexpected behaviors

# Trylock and Ordering

---

**Problem: Undesirably strong semantics for programs using non-blocking calls to acquire lock**

—e.g., `pthread_mutex_trylock()`

T1

`X=42;`

`lock(l);`

T2

`while(trylock(l)==success)`

`unlock(l);`

`assert(x==42);`

# The Problem with Trylock

---

T1	T2
X=42;	while(trylock(l)==success)
lock(l);	unlock(l);
	assert(x==42);

- **Problem: above program is data-race free by current semantics**
- **Assertion can fail if compiler or hardware reorders two statements executed by T1**
- **Prohibiting such reordering requires memory fence before lock**  
—fence doubles cost of the lock acquisition
- **For well-structured uses of locks and unlocks**  
—such reordering is safe and no fence is necessary



# C++ Memory Model Goals

---

- Sequential consistency for race free programs
- No semantics for programs with data races
  - different from Java      How?      Why?
- Weakened semantics for trylock

# Why Undefined Data Race Semantics?

---

- There are no benign data races
  - effectively the status quo
  - little to gain by allowing races other than allowing code to be obfuscated
- Giving Java-like semantics to data races may greatly increase cost of some C++ constructs
- Compiler optimizations often assume objects do not change unless there is an intervening assignment through a potential alias

```
unsigned i = x;  
if (i < 2) {  
    foo: ...  
    switch (i) {  
        case 0: ....; break;  
        case 1: ....; break;  
        default: ....;  
    }  
}
```

# Possible Memory Models

---

- **Sequential consistency**
  - intuitive, but restricts optimizations
- **Relaxed memory models**
  - allow hardware optimizations
  - specified at a low level making it hard to reason about correctness for programmers
  - limits compiler optimizations
- **Data-race free model**
  - properties
    - guarantee sequential consistency to data-race free programs
    - no guarantee for programs with races
  - simple model with high performance

# Definitions

---

- **Memory location**
  - each scalar value occupies a separate memory location
    - except bitfields inside the same innermost struct or class
- **Memory action consists of**
  - type of action
    - synchronization operation (for communication)
      - lock/unlock, atomic load/store, atomic read-modify-write
    - data operation: load, store
  - label identifying program point
  - values to be read and written

# Definitions

---

- **Thread Execution**
  - set of memory actions
  - partial order corresponding to the sequenced before ordering
    - sequenced before applies to memory operations by same thread
- **Sequentially Consistent Execution**
  - set of thread executions
  - total order,  $<T$ , on all the memory actions which satisfies the constraints
    - each thread is internally consistent
    - $T$  is consistent with sequenced-before orders
    - each load, lock, read-modify-write operation reads value from last preceding write to same locations according to  $<T$
- **Effectively requires total order that is interleaving of individual thread actions**

# Definitions Continued

---

- Two memory operations conflict if
  - they access same memory location
  - at least one of operation is a
    - store
    - atomic store
    - atomic read-modify-write
- Type 1 data race
  - in sequentially consistent operation, two memory operations from different threads form a type 1 data race if they conflict
    - at least one is a data operation
    - adjacent in  $<T$

# C++ Memory Model

---

- If a program (on a given input) has a sequentially consistent execution with a (type 1) data race, its behavior is undefined
- Otherwise, program behaves according to one of its sequentially consistent executions

# Legal Reorderings

---

**Hardware and compilers may freely reorder memory operation M1 sequenced before memory operation M2**

**if the reordering is allowed by intra-thread semantics and**

- 1. M1 is a data operation and M2 is a read synchronization operation**
- 2. M1 is a write synchronization and M2 is a data operation**
- 3. M1 and M2 are both data with no synchronization sequence-ordered between them**



# Legal Lock Optimizations

---

- When locks and unlocks are used in “well-structured” ways, following reorderings between M1 sequenced-before M2 are safe
  - M1 is data and M2 is the write of a lock operation
  - M1 is unlock and M2 is either a read or write of a lock
- Note: Data writes and writes from well-structured locks and unlocks can be executed non-atomically

# Trylock Revisited

---

- **Problem: trylock allows one to infer that a lock has been acquired**
- **Want to prevent this to avoid the cost of and additional fence**
- **Previous work achieved this by distinguishing between different types of synchronization operations and/or redefining data operations**
  - data-race-free-1 model distinguishes pairable and unpairable synchronization
  - Java has happens before

# C++ Memory Model Solution for Trylock

---

- **Modify specification of trylock to not guarantee that it will succeed if the lock is available**
- **Failed trylock doesn't tell you anything reliable**
  - can't use trylock to infer that another thread holds the lock
- **New semantics**
  - successful trylock is treated as lock()
  - unsuccessful trylock() is treated by memory model as no-op

# Implications for Current Processors

---

- Atomic writes need to be mapped to xchg on AMD64 and Intel64
- Hardware now only needs to make sure that xchg writes execute atomically
- xchg implicitly ensures semantics of a store|load fence
- Why is this OK?
  - better to pay a penalty on stores than on loads
  - store|load fence replaced by read-modify-write is as expensive on many processors today

# Atomic Ambiguity

---

- **Atomic can either be**
  - description of how a write is executed in hardware
  - C++ qualifier to denote special class of memory operations
- **Default C++ atomic operator needs to be executed atomically by hardware**

# Some Problems

---

- Significant restrictions on synchronization operations
- Synchronization operations must appear sequentially consistent with respect to each other
- Performance problem in practice
  - only Itanium distinguishes between data and sync operations
  - other processors enforce ordering through fence or memory barrier instructions
- Atomic operations must execute in sequenced-before and atomic writes must execute atomically
  - read can't return a new value for a location until all older copies of a location are invalid
  - with caches, atomic writes are easier with invalidation protocols

# Problematic Examples

---

- **Common to use counters that are frequently incremented but only read after all threads complete**
  - problem: requires all memory updates performed prior to counter update become visible after any later counter update in another thread
- **Atomic store requires two fences**
  - one before and one after the store
  - could imagine example where memory updates before or after store could be reordered

# Case for Low-Level Atomics

---

- **Cases exist where current model provides too much safety**
- **Expert programmers need way to relax model when code does not require as much safety to maximize performance**
- **Don't want to make memory model hard to reason about for the rest of us**



# Additional Definitions

---

- **Happens-before**
  - if a is sequenced before b, then a happens before b
  - if a synchronizes with b, then a happens before b
  - if a happens before b and b happens before c, then a happens before c
- **Type-2 data race**
  - two data conflicting accesses to the same memory location are unordered by happens before

# Modified C++ Memory Model

---

- If a program(on a given input) has a consistent execution with a (type 2) data race, then its behavior is undefined
- Otherwise, program(on the same input) behaves according to one of its sequentially consistent executions

# Properties

---

- **Type 2 data race**
  - if two data accesses to the same memory location are unordered by happens before
- **This model provides sequential consistency to programs whose consistent executions do not contain a type 2 data race**
- **If a program allows a type 2 data race in a consistent execution, then there exists a sequentially consistent execution, with two conflicting actions, neither of which happens before the other**
- **A program allows a type 2 data race of a given input iff there exists a sequentially consistent execution in which two unordered conflicting actions are adjacent in the sequential interleaving**

# Low-Level Atomics Support

---

- Memory model provides low-level atomic operations
- Expert programmer can maximize performance
- Atomic variables can be explicitly parameterized with respect to memory ordering constraint `x.load(memory_order_relaxed)` allows instruction to be reordered with other memory operations.
  - load is never an acquire operation, hence does not contribute to the synchronizes-with ordering
- For read-modify-write operations, programmer can specify whether the operations acts as an acquire, a release operation, neither (relaxed), or both

# Java/C++ Comparison

---

- **Primary goal of Java is safety and security**
  - make large effort to ensure semantics of codes
- **C++ focus is performance**
  - no such safety concerns
    - ignore semantics of codes with data races
  - low-level atomics to allow performance fine tuning

# Summary

---

- **Need to write parallel programs to get performance gains on modern architectures**
- **Previously, C++ memory model was ambiguous**
- **C++ memory model provides**
  - well defined semantics for data race free programs
  - high performance by leaving programs with data races undefined
  - low level atomics for expert programmers to squeeze out maximum performance
  - compilers may assume that ordinary variables don't change asynchronously
- **Remaining problems: adjacent bitfields**